

Netlink and Netlink for Wifi

Comunicações Móveis

2021/2022

DETI – UA

Daniel Corujo

Outline

- Sockets
- Netlink library (libnl)
- nl80211 library

Netlink Library (libnl) in a slide

- Allows the use of Netlink sockets Communication
 - Connecting/Disconnecting
 - Sending/Receiving Data
 - Message construction/parsing
 - Message reception state machines
 - Data structures
- Uses the Netlink Protocol
 - Socket-based inter-process Communications mechanism
 - User-space processes <-> kernel
 - User-space process <-> User-space process
 - Is a datagram-oriented service

(What is a socket?)

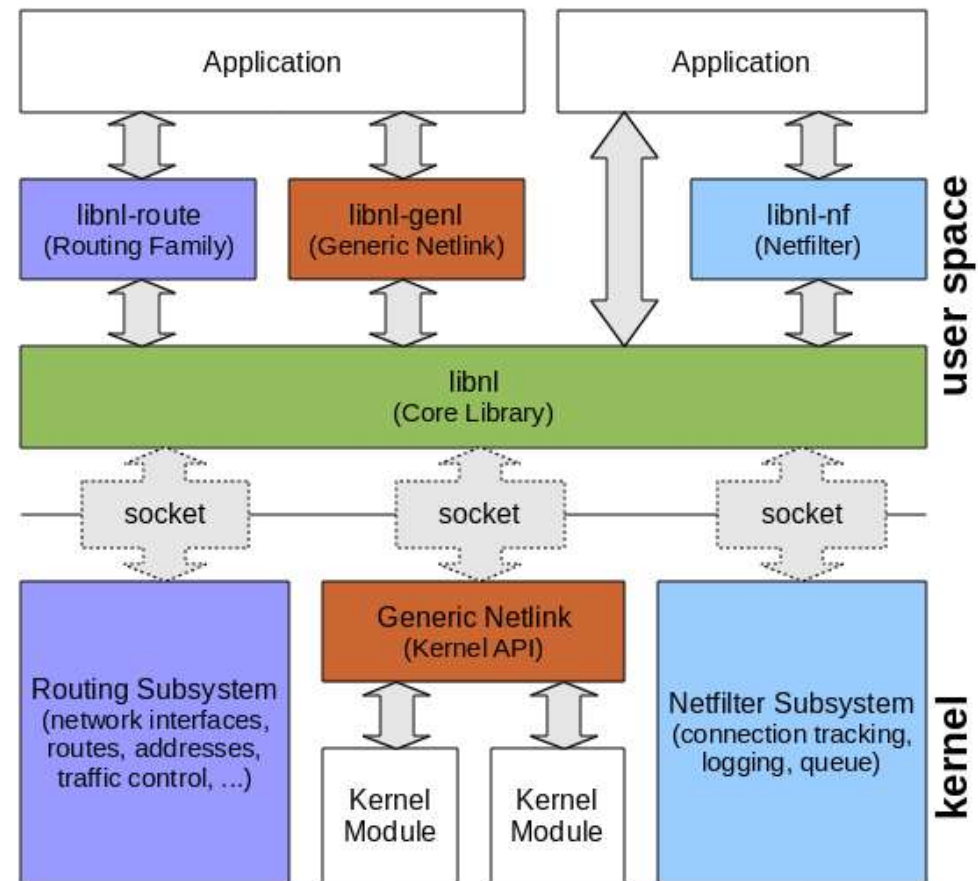
- Abstraction of an end-point for exchanging data between processes
 - Pipeline
 - `history | grep gcc`
 - Network Sockets
 - Between different hosts
 - Berkeley Sockets
 - Unix Domain socket / IPC socket
 - Internet Socket
 - Netlink Sockets
 - Host-only
 - Based on BSD sockets

(What is a socket)

- Berkeley socket uses a File Descriptor
 - Identifies an object for system resources
 - File descriptors cannot be directly accessed by user-space processes.
 - User-space processes does a system call to the kernel, providing it with the file descriptor reference, and the kernel accesses the resource (i.e., input/output) on its behalf
- Netlink socket used to use PID
 - Not anymore (multi-threading!)
 - Now uses a 32-bit port number

Netlink library (libnl)

- Libnl
 - Netlink Library
 - Socket handling
 - Sending and receiving
 - Message construction and parsing
- Libnl-genl
 - Generic Netlink Library
 - Controller API
 - Family and command registration



Netlink socket families

- AF_NETLINK
 - Supports different subsects, each targeting different kernel components and messages (i.e., address families)
- Socket types
 - SOCK_RAW
 - SOCK_DGRAM
- Netlink family
 - Selects the kernel module or netlink group to communicate with

Netlink Families


- NETLINK_ROUTE
- NETLINK_W1
 - 1-wire
- NETLINK_USERSOCK
- NETLINK_FIREWALL
- NETLINK_SOCK_DIAG
- NETLINK_NFLOG
 - Netfilter/iptables log
- NETLINK_XFRM
- NETLINK_SELINUX
- NETLINK_ISCSI
- NETLINK_AUDIT
- NETLINK_FIB_LOOKUP
- NETLINK_CONNECTOR
- NETLINK_NETFILTER
- NETLINK_GENERIC
- ...

Netlink Protocol

Netlink Protocol

- Socket-based inter-process Communications mechanism
 - User-space processes <-> kernel
 - User-space process <-> User-space process
- Is a datagram-oriented service

Datagram Message Format




- Message Type
 - Type of payload
- Message Flags
 - Modify the behavior of the message type
 - Request
 - Multicast
 - Acknowledgement
 - Echo

Types of Netlink Messages


- Requests
- Notifications
- Replies

Message Types (datagram parameter)

- NLMSG_NOOP
 - No operation, message must be discarded
- NLMSG_ERROR
 - Error message or ACK
- NLMSG_DONE
 - End of multipart sequence
- NLMSG_OVERRUN
 - Overrun notification (Error)



Message Flags



- NLM_F_REQUEST - Message is a request
- NLM_F_MULTI - Multipart message
- NLM_F_ACK - ACK message requested
- NLM_F_ECHO - Request to echo the request
- Universal flags for GET requests
 - NLM_F_ROOT - Return based on root of tree.
 - NLM_F_MATCH - Return all matching entries.
 - NLM_F_ATOMIC - Obsoleted, once used to request an atomic operation.
 - NLM_F_DUMP - Return a list of all objects (NLM_F_ROOT|NLM_F_MATCH).

Netlink Sockets

Netlink sockets

- Needed to use the Netlink protocol
- It's where you send and receive protocol messages

Program example

- Let's build a program that waits for NETLINK_ROUTE notifications. When one is received, it calls a user function (i.e., **callback function**)
- We need to:
 - Create a socket
 - Indicate which is the callback function
 - Connect the socket to NETLINK_ROUTE
 - Subscribe for receiving specific notifications, from a specific multicast group
 - Keep the program running, listening for the notifications
 - *(a notification is an event that is sent when some specific action occurred)*

Program example

//this program waits for NETLINK_ROUTE notifications. When one is received, it calls a user function (i.e., **callback function**)

//function

```
static int my_func(struct nl_msg *msg, void *arg){  
    printf("A message was received! This is my function");  
    return 0; }
```

Program example

```
//Allocate a new socket
```

```
stuct nl_sock * sk = nl_socket_alloc();
```

```
//Disable sequence number checking, as we're just using notifications
```


```
//Netlink automatically takes care of sequence numbers, when using  
the nl_send_auto() for sending messages.
```

```
//However, if we're using a non request/reply netlink protocol, we must  
explicitly disable it:
```

```
nl_socket_disable_seq_check(sk);
```

Program example

//Callback function that calls “my_func”



```
nl_socket_modify_cb(sk, NL_CB_VALID, NL_CB_CUSTOM, my_func,  
null);
```

//Connect to routing netlink protocol

```
nl_connect(sk, NETLINK_ROUTE);
```

```
int nl_socket_modify_cb(struct nl_sock *sk, enum nl_cb_type type, enum nl_cb_kind kind,  
                        nl_recvmsg_msg_cb_t func, void *arg);
```

Program example

```
//subscribe to link notifications multicast group  
nl_socket_add_membership(sk, RTNLGRP_LINK,0);
```

```
//Listen to messages
```

```
//This will trigger the callback function, when a notification of  
NETLINK_ROUTE protocol is received
```

```
while (1)  
    nl_rcvmsgs_default(sk);
```

Socket related functions


- `nl_socket_set_local_port`
- `nl_socket_set_peer_port`
- `nl_socket_get_fd`
- `nl_socket_set_buffer_size`
- ...

Sending messages through a netlink socket

- Two ways
 - Automatic
 - `nl_send_auto()`
 - Automatically completes the missing bits of the netlink message header
 - Automatically deals with addressing based on current information set in the socket
 - Passes the message to `nl_send()`
 - Creates the final message
 - Manual (when automatic filling is not suitable)
 - Directly use `nl_send()`
 - Embeds the message into a 'iovec' structure and pass it to `nl_send_iovec()`
 - `nl_send()` can also be overwritten via `nl_cb_overwrite_send()`

Sending messages through a netlink socket

- Both ways lead to
 - `nl_send_iovec()`
 - Fills the message header – tries to find the peer, or leaves it to the kernel
 - Leads to:
 - `nl_sendmsg()`
 - Takes the final message and the optional header
 - Sends the final total message to `sendmsg()`



Receiving messages through a netlink socket

- Our program example received messages (Notifications) from the kernel using the following function:

```
while (1)
    nl_recvmsgs_default(sk);
```

- Easiest function
 - Receives messages based on how we configured the socket
 - Usually, the default behavior is enough
 - Fetches the callbackfunction (`cb = nl_socket_get_cb(sk)`)
 - calls `nl_recvmsgs()`

Receiving messages through a netlink socket

- `nl_recvmsgs()`
 - Actual message reception loop
 - If we need specific reception characteristics, we can provide a complete own implementation of the reception mechanism
 - `nl_cb_overwrite_recvmsgs()`

Parsing messages

- Messages are 4-bytes aligned in all boundaries
- There are two methods of parsing
 - Low-level interface (manual parsing)
 - High-level interface (Implement a parser as part of cache operations)
- What is receiving a netlink protocol message on a netlink socket?
 - What you receive from a netlink socket is typically a stream of messages.
 - You will be given a buffer and its length
 - The buffer may contain any number of netlink messages.
 - The first message header starts at the beginning of the message stream
 - You can reach the next header by calling `nlmsg_next()` on the previous header
 - $\text{Position} = \text{Remaining_number_of_bytes} - \text{current_message_size}$

Parsing messages

- Despite having `nlmsg_next()` we don't know if there are more messages
- We must assume that more messages follow until all bytes of the stream have been processed
 - `nlmsg_ok()`
 - Returns true if another message fits into the remaining number of bytes in the message stream
 - `nlmsg_valid_hdr()`
 - Checks if a message contains at least a minimum of payload

Creating a parsing message function: example

```
#include <netlink/msg.h>

void my_parse(void *stream, int length){

    struct nlmsghdr *hdr = stream;

    while (nlmsg_ok(hdr, length)) {

        // Parse message here

        hdr = nlmsg_next(hdr, &length);
    }
}
```

Parsing messages

- Accessing message payload
 - Remember that the header has alignment
 - Some of its fields, and the header itself, might have padding
 - `nlmsg_data()` returns a pointer to the start of the payload
 - `nlmsg_dataalen()` returns the length of the message payload
 - `nlmsg_tail()` return a pointer to the end of the payload, including padding



Attributes

- Most netlink protocol messages use netlink attributes
- This means that, the message payload is composed of
 - Protocol Header (+ padding)
 - Attributes (+ padding)
- `nlmsg_attrdata()` returns a pointer to the beginning of the attributes section
- `nlmsg_attrlen()` returns the length of the attributes section

Parsing function

```
int nlmsg_parse(struct nlmsg_hdr *hdr, int hdrlen, struct nlattr **attrs,  
               int maxtype, struct nla_policy *policy);
```

- nlmsg_parse()
 - Starts by validating the header
 - If hdrlen>0, calls nlmsg_valid_hdr()
 - Fills an array with pointers to each attribute
- nlmsg_validate()
 - Similar, but does not create the array
- There are also attribute variants of this function
 - nla_parse()
 - nla_validate() → does not create the array

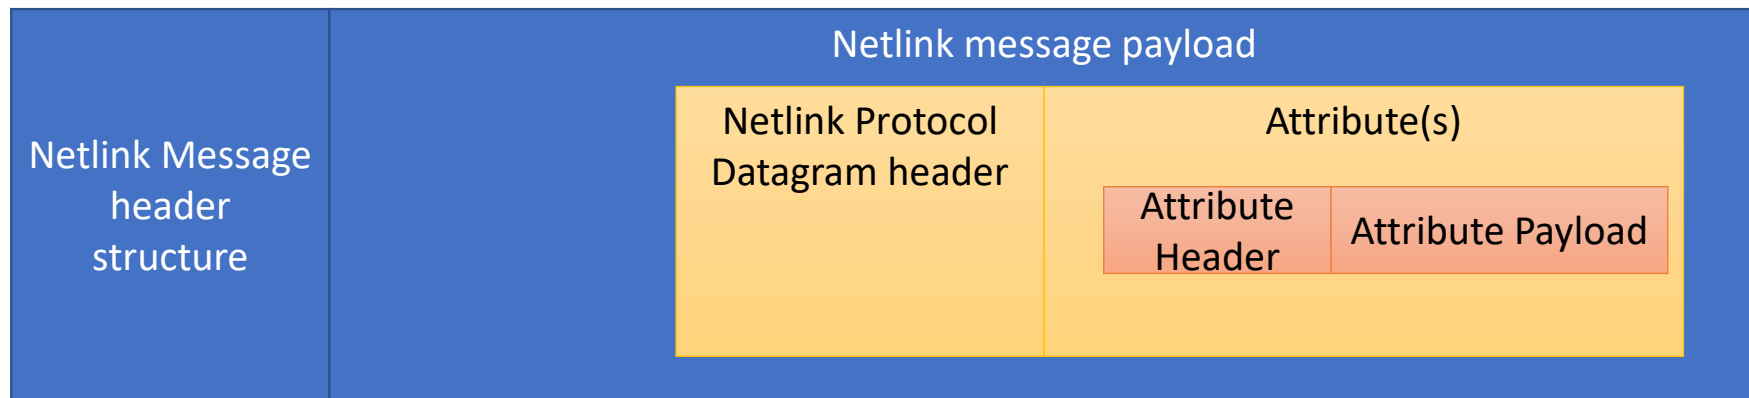
Creating a message

- We need to allocate a struct `nl_message`
 - Holds the message header and payload
- `nlmsg_alloc()` is the function used for message allocation
 - Uses the default maximum message size (one page, typically 4K)
 - Default size can be changed with `nlmsg_set_default_size(size_t)`
 - `nlmsg_alloc_size()` is a variant that allows definition of size at allocation time
 - Alternative #1: If we want to reuse an already known header, we can call `nlmsg_inherit()`
 - Appends header
 - Alternative #2: `nlmsg_alloc_simple()` takes a message type and flag
 - Creates and appends header

Create a message

- If we don't automatically create a header (`nlmsg_alloc_simple()` or `nlmsg_inherit()`) we need to add the header ourselves
- `nlmsg_put()`
 - `nlmsg_type`
 - `nlmsg_flags`
 - `seqnr (NL_AUTO_SEQ)`
 - `Port (NL_AUTO_PORT)`

Do not get confused!



Example

```
#include <netlink/msg.h>
```

```
struct nlmsghdr *hdr; //header structure
```

```
struct nl_msg *msg; //message structure
```

```
struct myhdr {
```

```
    uint32_t foo1, foo2;
```

```
} hdr = { 10, 20 }; //my header
```

Example (contd)

```
msg = nlmsg_alloc(); //allocate default maximum size message

// Create header with automatic fill of port and sequence number
//Notice that space for hdr is reserved in the message
hdr = nlmsg_put(msg, NL_AUTO_PORT, NL_AUTO_SEQ, MY_MSGTYPE,
sizeof(hdr), NLM_F_CREATE);

//copy customised header into payload
memcpy(nlmsg_data(hdr), &hdr, sizeof(hdr));
```

What about the payload?

- Payload should be encoded as netlink attributes whenever possible
- Attributes are aligned at multiples of 4bytes position
 - `nla_padlen()` returns the number of padding bytes needed
- Attribute encoding
 - Length (16bit)
 - Type (16bit)
 - Payload

Parsing attributes

- You can obtain them as an array, when you parse the message with `nlmsg_parse()`
- We can also do it manually
 - Navigate individual attributes using `nla_next()` after getting a pointer to the first one using `nlmsg_attrdata()`
 - `nla_ok()` determines whether another attribute fits into the remaining number of bytes or not

Parsing Attributes

- Each attribute has a header and payload
- Access the header: `nla_len()`
- Access the type: `nla_type()`
- Access the payload: `nla_data()`
 - Avoid casting to anything larger than 4bytes due to the mandatory alignment
- Easier alternative: `nla_parse()`
 - Iterates all attributes in the stream
 - Validates each attribute
 - When validation succeeds, stores pointers to attributes in array

Attribute Validation

- We need to ensure that we get attributes formatted in the correct way
- We define structures that compose policies
 - Indicate what is the structure of the attribute's header

```
struct nla_policy {  
    uint16_t type;    //NLA_U32/16/8/4  
    uint16_t minlen;  //minimum payload length  
    uint16_t maxlen; }; //maximum payload length
```

S

Attribute parsing example

//headers are ignored in this example

```
#include <netlink/msg.h>
```

```
#include <netlink/attr.h>
```

```
enum {
```

```
    MY_ATTR_FOO = 1, MY_ATTR_BAR, __MY_ATTR_MAX,};
```

```
#define MY_ATTR_MAX (__MY_ATTR_MAX - 1)
```

Attribute parsing example (contd)

```
static struct nla_policy my_policy[MY_ATTR_MAX+1] = {  
    //Validation policy  
    [MY_ATTR_FOO] = { .type = NLA_U32 },  
  
    [MY_ATTR_BAR] = { .type = NLA_STRING,  
  
        .maxlen = 16 },  
  
};
```

Attribute parsing example (contd)

```
void parse_msg(struct nlmsghdr *nlh)

    struct nlattr *attrs[MY_ATTR_MAX+1];

    if (nlmsg_parse(nlh, 0, attrs, MY_ATTR_MAX, my_policy) < 0) /* error */

    if (attrs[MY_ATTR_FOO]) {
        /* MY_ATTR_FOO is present in message */
        printf("value: %u\n", nla_get_u32(attrs[MY_ATTR_FOO]));
    }
```

I want to find a single attribute!

- There are functions that iterate over all attributes, search for a matching one and return a pointer to its header
- `nla_find()`;
- `nlmsg_find_attr`

Nested attributes

- Attributes included as payload of container attributes
 - Type: NLA_NESTED
- Attributes can be stored inside a tree structure
- It is the common way to transmit lists of objects

Parsing Nested Attributes

- `nla_parse_nested()`
 - Identical to `nla_parse()`, but
 - Uses a struct `nlattr` as argument
 - Uses the payload as stream of attributes

Constructing Nested Attributes

- We can nest attributes by surrounding them with
 - `nla_nest_start()`
 - Add attribute header, without payload
 - All data added from this point on, will be part of the container
 - `nla_nest_end()`
 - Closes the container attribute

Create Netlink Message with Attributes

//We are going to put this behavior inside of a function

//So everything you'll see next will be placed inside this function

```
struct nl_msg *build_msg(int ifindex, struct nl_addr *lladdr, int mtu)
{
    //Code from the next slides will be here
}
```

Create Netlink Message with Attributes (cont)

```
struct nl_msg *msg;
```

```
struct nlattr *info, *vlan;
```

```
struct ifinfomsg ifi = {
```

```
    .ifi_family = AF_INET,
```

```
    .ifi_index = ifindex,
```

```
};
```

Estrutura predefinida no netlink:

```
struct ifinfomsg {  
    unsigned char ifi_family; /* AF_UNSPEC */  
    unsigned short ifi_type; /* Device type */  
    int ifi_index; /* Interface index */  
    unsigned int ifi_flags; /* Device flags */  
    unsigned int ifi_change; /* change mask */ };
```

Create Netlink Message with Attributes (cont)

```
/* Allocate a default sized netlink message */
```

```
//ROUTING FAMILY NETLINK
```

```
if (!(msg = nlmsg_alloc_simple(RTM_SETLINK, 0)))
```

```
    return NULL;
```

Create Netlink Message with Attributes (cont)

```
/* Append the protocol specific header (struct ifinfomsg)*/
```

```
if (nlmsg_append(msg, &ifi, sizeof(ifi), NLMSG_ALIGNTO) < 0)
```

```
goto nla_put_failure;
```

Create Netlink Message with Attributes (cont)

`/* Append a 32 bit integer attribute to carry the MTU */`

```
NLA_PUT_U32(msg, IFLA_MTU, mtu);
```

`/* Append a unspecific attribute to carry the link layer address */`

```
NLA_PUT_ADDR(msg, IFLA_ADDRESS, lladdr);
```

Create Netlink Message with Attributes (cont)

```
/* Append a container for nested attributes to carry link information */
```

```
if (!(info = nla_nest_start(msg, IFLA_LINKINFO)))
```

```
    goto nla_put_failure;
```

```
/* Put a string attribute into the container */
```

```
NLA_PUT_STRING(msg, IFLA_INFO_KIND, "vlan");
```

Create Netlink Message with Attributes (cont)

```
/* Append another container inside the open container to carry  
 * vlan specific attributes */
```

```
if (!(vlan = nla_nest_start(msg, IFLA_INFO_DATA)))
```

```
    goto nla_put_failure;
```

```
/* add vlan specific info attributes here... */
```

Create Netlink Message with Attributes (cont)

```
/* Finish nesting the vlan attributes and close the second container. */
```

```
    nla_nest_end(msg, vlan);
```

```
/* Finish nesting the link info attribute and close the first container. */
```

```
    nla_nest_end(msg, info);
```

```
return msg;
```


Create Netlink Message with Attributes (cont)

```
//code alias for failures:
```

```
nla_put_failure:
```

```
    nlmsg_free(msg);
```

```
    return NULL;
```

```
}
```

Parsing a Netlink Message with Attributes

//We are going to put this behavior inside of a function

//So everything you'll see next will be placed inside this function

```
int parse_message(struct nlmsghdr *hdr)
```

```
{
```

```
    // The code from the next slides will be here
```

```
}
```

Parsing a Netlink Message with Attributes

```
/* The policy defines two attributes: a 32 bit integer and a container  
for nested attributes.      */
```

```
struct nla_policy attr_policy[] = {  
    [ATTR_FOO] = { .type = NLA_U32 },  
    [ATTR_BAR] = { .type = NLA_NESTED },    };
```

```
struct nlattr *attrs[ATTR_MAX+1];  
int err;
```

Parsing a Netlink Message with Attributes

```
/* The nlmsg_parse() function will make sure that the message  
contains enough payload to hold the header (struct my_hdr), validates  
any attributes attached to the messages and stores a pointer to each  
attribute in the attrs[] array accessible by attribute type.    */
```

```
if ((err = nlmsg_parse(hdr, sizeof(struct my_hdr), attrs, ATTR_MAX,  
attr_policy)) < 0)  
    goto errout;
```

```
if (attrs[ATTR_FOO]) {
```

```
    /*It is safe to directly access the attribute payload without any  
    further checks since nlmsg_parse() enforced the policy. */
```

```
    uint32_t foo = nla_get_u32(attrs[ATTR_FOO]);
```

```
}
```

```
if (attrs[ATTR_BAR]) {  
    struct *nested[NESTED_MAX+1];  
    /* Attributes nested in a container can be parsed the same way as top level attributes. */  
  
    err = nla_parse_nested(nested, NESTED_MAX, attrs[ATTR_BAR], nested_policy);  
  
    if (err < 0)  
        goto errout;  
  
    // Process nested attributes here.  
  
}
```

More details

- Core Library Developer's Guide
 - <https://www.infradead.org/~tgr/libnl/doc/core.html>
- API Reference
 - https://www.infradead.org/~tgr/libnl/doc/api/group_core.html

Genery Netlink Library

Libn-genl

- One of the drawbacks of the Netlink protocol is that the number of protocol families is limited to 32 (MAX_LINKS)
- For this reason, the “Generic Netlink Family” was created
 - Provides support for adding a higher number of families
 - Acts as a Netlink multiplexer
 - Works with a single Netlink family: NETLINK_GENERIC
 - The generic Netlink protocol is based on the Netlink protocol and used its API

Generic Netlink Library – Controller API

- The controller is a component in the kernel that resolves Generic Netlink family names to their numeric identifiers. This module provides functions to query the controller to access the resolving functionality.
- Example:
 - `expectedId = genl_ctrl_resolve(sk, "nl80211");`

Generic Netlink Library – Controller API

- Genlmsg_put adds Generic Netlink headers to Netlink messages
 - `genlmsg_put(msg, 0, 0, driver_id, 0, NLM_F_DUMP, NL80211_CMD_GET_SCAN, 0);` // Setup which command to run.

Generic Netlink Library – Controller API

- Family and command registration
- Registers the specified Generic Netlink family definition together with all associated commands. After registration, received Generic Netlink messages can be passed to **genl handle msg()** which will validate the messages, look for a matching command and call the respective callback function automatically.

How can I use lib-nl and lib-genl

- You need to install the libs
 - libnl-3-dev
 - libnl-genl-3-dev
- The code you do will have to include related headers
 - `#include "netlink/netlink.h"`
 - `#include "netlink/genl/genl.h"`
 - `#include "netlink/genl/ctrl.h"`

To compile

- `gcc example.c -o example -I/usr/include/libnl3/ -lnl-genl-3 -lnl-3`
- The `-I` holds the location of Include header files, which are referenced in the `#include` directive. E.g. `-I /home/myname/include` will search for a header file in `"/home/myname/include"`.
- Since in Unix/Linux OSes library names begin with `"lib"` (e.g `"libkuku.a"`), then the `"lib"` prefix is omitted from the `-l` parameter: `"-lkuku"`.
- (Not used here, but the `-L` holds the location of library files, and `-l` holds the name of a library file, which should be included in the Link phase of the program's build.)

What has this to do with Wi-Fi?

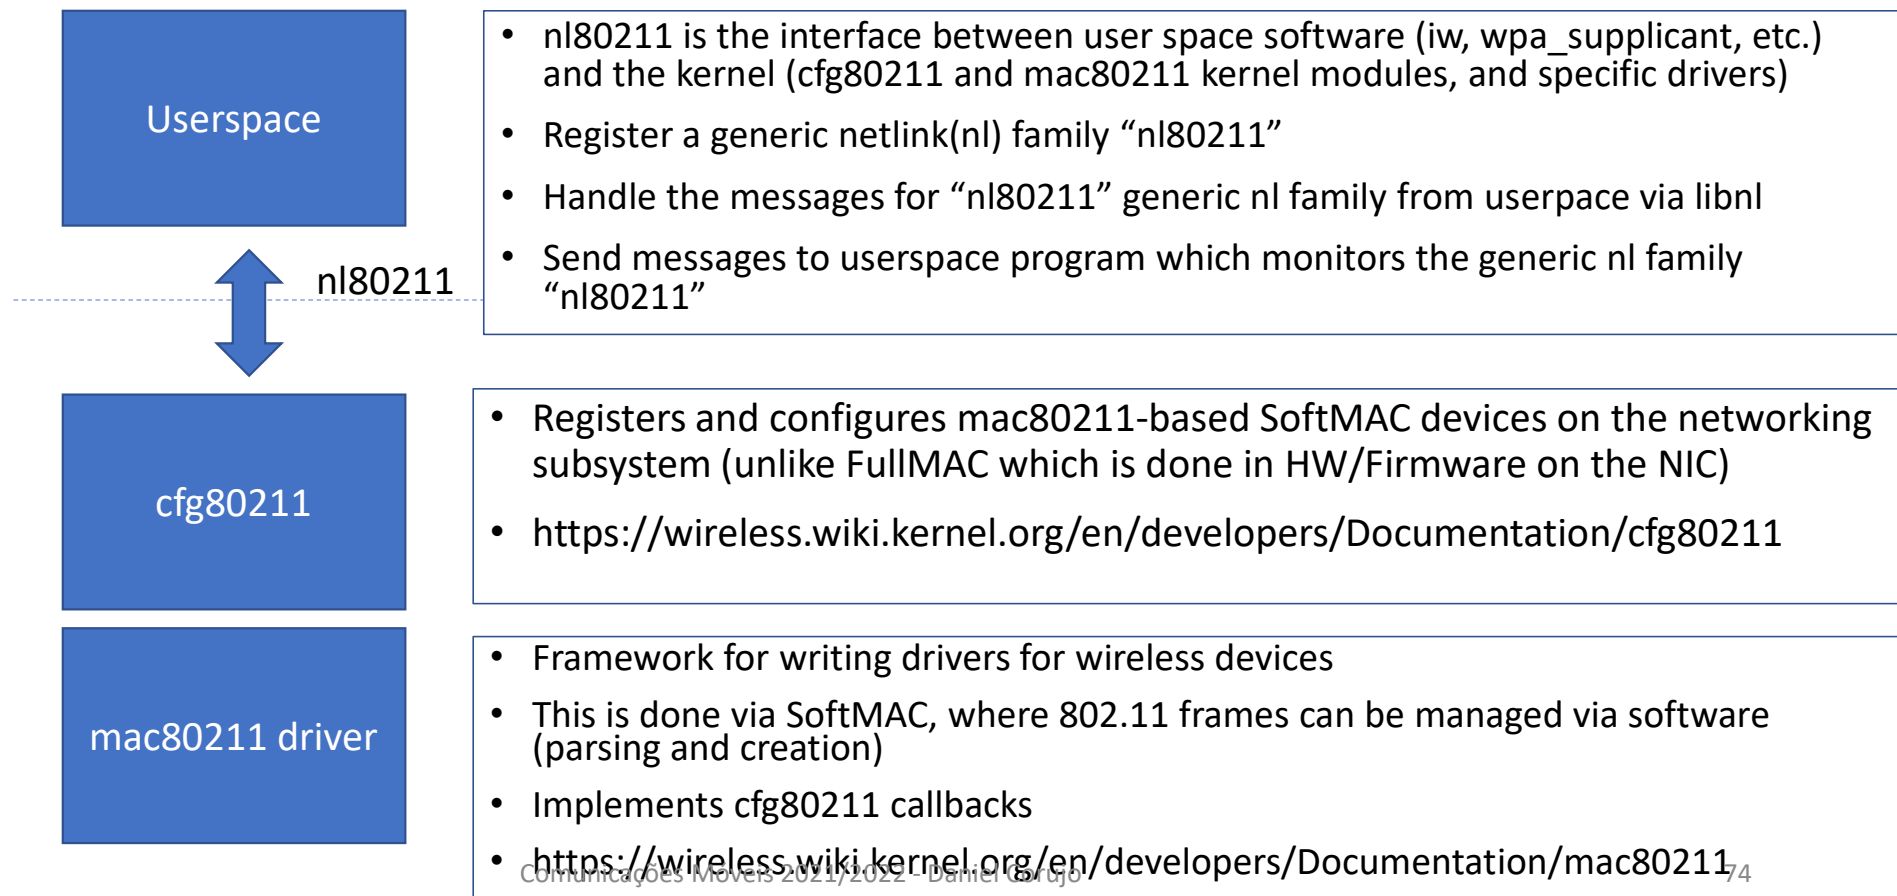
- Until now, we have been addressing how the netlink library works
- We will now address a Wi-Fi related family...

NL80211

NL80211

- 802.11 userspace interface for Netlink
- Allows control of wireless drivers from userspace, using IPC Communications between kernel and userpace
- Before:
 - Input/output control (ioctl) system calls
 - Wireless Extensions generic API

NL80211



Can I use this?

- Only works on network interfaces whose drivers are compatible with Netlink.
- Test this by running ``iw list``
- Some of the commands used in the code require 'root' privileges
- The library's header files already come installed with the kernel

Examples of NL80211-based apps

- iw
 - Wireless devices CLI configuration tool
- crda
 - Helper for regulatory domains compliance using udev device management
 - Country code, calibration channel, calibration bandwidth
- hostapd
 - Access Point connection and authentication daemon
- wpa_supplicant (-Dnl80211)
- iwd
 - Wireless management service
- NetworkManager
 - Connection manager based on iwd
- ConnMan
 - Another Connection Manager (CLI)

Role of nl80211.h

- Provides several enums related with 802.11
- You then use generic nl80211 constructs (i.e., sockets, messages, attributes, commands) to send commands to the wireless card

Generic structure of a nl80211 program

1 - Allocate new netlink socket in memory

```
struct nl_sock* sk = nl_socket_alloc();
```

2 – Create file descriptor and create socket

```
genl_connect(sk);
```

3 - Find the nl80211 driver ID.

```
int driver_id = genl_ctrl_resolve(socket, "nl80211");
```

Generic structure of a nl80211 program

- From here, you can do whatever you want
- Typically:
 - Create a callback function and attach it to the socket
 - Create a message with a command to be sent to the wifi card
 - Header + attribute(s)
 - Send the message and wait for answer
 - Might involve joining some multicast group
 - When reply arrives, your callback function is called
 - Inside of that function, you build code that processes the message payload and attributes (i.e., prints information in screen)

Create a callback function and attach it to the socket

- `nl_socket_modify_cb(sk, NL_CB_VALID, NL_CB_CUSTOM, yourFunctionName, NULL);`

Create a message with a command to be sent to the wifi card

```
//allocate a message
struct nl_msg* msg = nlmsg_alloc();

//request interface's configuration
enum nl80211_commands* cmd = NL80211_CMD_GET_INTERFACE;
int ifIndex = if_nametoindex("wlan0");
int flags = 0;

// setup the message
genlmsg_put(msg, 0, 0, driver_id, 0, flags, cmd, 0);

//add message attributes
NLA_PUT_U32(msg, NL80211_ATTR_IFINDEX, ifIndex);
```

Send the message and wait for answer

//send the message (this frees it)

```
ret = nl_send_auto_complete(sk, msg);
```

//block for message to return

```
nl_recvmsgs_default(sk);
```

Callback function

```
static int nlCallback(struct nl_msg* msg, void* arg) {
    struct nlmsg_hdr* ret_hdr = nlmsg_hdr(msg);
    struct nlattr *tb_msg[NL80211_ATTR_MAX + 1];

    if (ret_hdr->nlmsg_type != expectedId) return NL_STOP;

    struct genlmsg_hdr *gnlh = (struct genlmsg_hdr*) nlmsg_data(ret_hdr);
    nla_parse(tb_msg, NL80211_ATTR_MAX, genlmsg_attrdata(gnlh, 0), genlmsg_attrlen(gnlh, 0), NULL);

    if (tb_msg[NL80211_ATTR_IFTYPE]) {
        int type = nla_get_u32(tb_msg[NL80211_ATTR_IFTYPE]);
        printf("Type: %d", type);
    }
}
```

Joining a multicast group

```
int mcid = nl_get_multicast_id(socket, "nl80211", "scan");  
nl_socket_add_membership(socket, mcid);
```

Executing a scan

- Use the “NL80211_CMD_TRIGGER_SCAN” command to start a scan
 - If you try to start another one when one is running, this would fail
- Listen for the scan to complete when you get a NL80211_CMD_NEW_SCAN_RESULTS
- You can then send a NL80211_CMD_GET_SCAN command to ask for the results
 - You will get one message back for every station found, so be ready to handle multiple messages

Examples of commands

- Detailed explanation: nl80211.h
- Commands
 - NL80211_CMD_GET_INTERFACE
 - Request na interface's configuration
 - For all interfaces or a specific one
 - For a single interface it is necessary to send the request with the attribute NL80211_ATTR_IFINDEX
 - NL80211_CMD_SET_BEACON
 - Change the beacon on na access point interface using attributes
 - NL80211_CMD_GET_STATION
 - Get attributes for station identified by NL80211_ATTR_MAC on the interface identified by NL80211_ATTR_IFINDEX
 - NL80211_CMD_GET_REG
 - Obtain regulatory domain information

More information

- Unfortunately, there is no developer's guide for nl80211
- The best source of (properly coded) information, is to look at the reference implementations (iw)
- Some simpler code snips may be found in individual projects online
 - <https://github.com/Robpol86/libnl>